

Core Stack Switching Implementation Strategies (discussion)

Daniel Hillerström

The University of Edinburgh, UK

May 19, 2025

WebAssembly Stacks Subgroup

<https://dhil.net>

Implementation strategies

- Stack as a linked list of frames (e.g. Danvy (1987))
- Stack copying and pasting (e.g. Leijen (2017))
- Virtual memory mapped stacks (e.g. Leijen and Sivaramakrishnan (2021))
- Segmented stacks (e.g. Bruggeman, Waddell, and Dybvig (1996), Alvarez-Picallo et al. (2024))
- Whole-program transformations (e.g. Hillerström (2022), Zakai (2019))

Implementation strategies

- Stack as a linked list of frames (e.g. Danvy (1987))
- **Stack copying and pasting** (e.g. Leijen (2017))
- **Virtual memory mapped stacks** (e.g. Leijen and Sivaramakrishnan (2021))
- **Segmented stacks** (e.g. Bruggeman, Waddell, and Dybvig (1996), Alvarez-Picallo et al. (2024))
- Whole-program transformations (e.g. Hillerström (2022), Zakai (2019))

Methodology

Setup

- Source language: C with a bespoke fiber library
 - Stack copying: Libhandler backend (Leijen 2017)
 - Virtual memory stacks: Libmprompt backend (Leijen and Sivaramakrishnan 2021)
 - Segmented stacks: Libseff backend (Alvarez-Picallo et al. 2024)
- Disclaimer: mostly off-the-shelf configurations
- Three CPU-bound microbenchmarks:
 - Iterative summing
 - Recursive tree summing
 - Prime sieve computation
- Disclaimer: not particularly representable of ‘real-world’ workloads
- Using hyperfine to compare benchmarks end-to-end

Source code: github.com/dhil/core-stack-switching-experiments

Fibers interface in C

```
/** The signature of a fiber entry point. */
typedef void* (*fiber_entry_point_t)(void*);
/** The abstract type of a fiber object. */
typedef struct fiber* fiber_t;

/** Allocates a new fiber with the default stack size. */
fiber_t fiber_alloc(fiber_entry_point_t entry);
/** Reclaims the memory occupied by a fiber object. */
void fiber_free(fiber_t fiber);

/** Yields control to its parent context. */
void* fiber_yield(void *arg);

/** Possible status codes for 'fiber_resume'. */
typedef enum { FIBER_OK, FIBER_YIELD, FIBER_ERROR } fiber_result_t;

/** Resumes a given 'fiber' with argument 'arg'. */
void* fiber_resume(fiber_t fiber, void *arg, fiber_result_t *result);
```

Iterative sum (1)

Characteristics

- 2 coroutines, repeatedly yielding control in a tight loop
- Shallow call stack depth

Benchmark essence

```
for (size_t i = 0; i < N; i++) {  
    result += fiber_yield(&i);  
}
```

Iterative sum (2)

		Implementations		
		Stack copying	Virtual stacks	Segmented stacks (baseline)
N	10000	4.11	1.71	1.0
	100000	8.85	2.45	1.0
	1000000	11.15	2.86	1.0

Lower is better

Tree sum (1)

Characteristics

- 2 coroutines
 - 1 coroutine walks the tree recursively
 - 1 coroutine yields in a tight loop
- Deep call stack depth

Benchmark essence

```
void walk_tree(node_t *node) {  
    if (node->tag == LEAF) {  
        fiber_yield(node->val);  
    } else {  
        walk_tree(node->left);  
        walk_tree(node->right);  
    }  
}
```


Tree sum (2)

		Implementations		
		Stack copying	Virtual stacks	Segmented stacks (baseline)
N	21	11.22	2.81	1.0
	22	11.32	2.83	1.0
	23	11.59	2.86	1.0

Lower is better

Prime sieve (1)

Characteristics

- Dynamically spawns coroutines (up to N)
- Each coroutine holds a unique prime number
- Shallow call stack depth

Benchmark essence

```
bool filter(int32_t my_prime) {  
    bool divisible = false;  
    int32_t candidate = fiber_yield(NULL);  
    while (candidate > 0) {  
        divisible = (candidate % my_prime) == 0;  
        candidate = fiber_yield(divisible);  
    }  
    return false; // dummy value  
}
```

Prime sieve (2)

		Implementations		
		Stack copying	Virtual stacks	Segmented stacks (baseline)
N	1000	9.36	3.87	1.0
	4000	8.09	2.87	1.0
	8000	7.79	2.65	1.0

Lower is better

References I

- Danvy, Olivier (1987). “Memory allocation and higher-order functions”. In: *PLDI*. ACM, pp. 241–252. doi: 10.1145/29650.29676. URL: <https://doi.org/10.1145/29650.29676>.
- Bruggeman, Carl, Oscar Waddell, and R. Kent Dybvig (1996). “Representing Control in the Presence of One-Shot Continuations”. In: *PLDI*. ACM, pp. 99–107. doi: 10.1145/231379.231395. URL: <https://doi.org/10.1145/231379.231395>.
- Leijen, Daan (2017). “Implementing Algebraic Effects in C - “Monads for Free in C””. In: *APLAS*. Vol. 10695. Lecture Notes in Computer Science. Springer, pp. 339–363. doi: 10.1007/978-3-319-71237-6_17. URL: https://doi.org/10.1007/978-3-319-71237-6_17.
- Zakai, Alon (2019). *Pause and Resume WebAssembly with Binaryen’s Asyncify*. Accessed 2022-10-27. URL: <https://kripken.github.io/blog/wasm/2019/07/16/asyncify.html>.
- Leijen, Daan and KC Sivaramakrishnan (2021). *libmprompt: Robust multi-prompt delimited control and effect handlers in C/C++*. <https://github.com/koka-lang/libmprompt>.
- Hillerström, Daniel (2022). “Foundations for Programming and Implementing Effect Handlers”. PhD thesis. The University of Edinburgh, UK. doi: 10.7488/era/2122. URL: <http://dx.doi.org/10.7488/era/2122>.

References II

Alvarez-Picallo, Mario et al. (2024). “Effect Handlers for C via Coroutines”. In: *Proc. ACM Program. Lang.* 8.OOPSLA2, pp. 2462–2489. DOI: 10.1145/3689798. URL: <https://doi.org/10.1145/3689798>.